

Operating Systems

CSCI 5806

Spring Semester 2025 — CRN 22968

Term Project — Step 1 — VDI File Access

Target completion date: Monday, January 27, 2025

Goals

- Provide the five basic file I/O functions to access disk space inside a VDI file.
- Create a structure or class to contain the data necessary to implement the five functions.

Details

You'll want a single entity — a structure or a class, either works — to represent a VDI file within your project. The intent is to collect the various data your project is going to need into one place for ease of use; if you're using C++ (you *are* using C or C++, right?) then a class can also contain the basic I/O functions as methods.

What do you need to keep track of in a VDI file? To implement the five functions you'll probably want at least these four items to be contained in your structure:

- The file descriptor for the VDI file. This is an `int`.
- A VDI header structure. Discussion of the structure can be found at <https://forums.virtualbox.org/viewtopic.php?t=8046>; C structures for the header can be found in the VirtualBox source code.
- The VDI translation map. This is optional; include it if you are going to enable dynamic VDI files (they're easy!), ignore it otherwise. This is just an array of integers, although you don't know the size in advance. The size is given in the VDI header, so you'll need to allocate this dynamically.
- A cursor. This is just an integer (of type `size_t`) that holds the location of the next byte to be read or written.

Once the structure / class is created, you'll need to implement a VDI version of each of the five basic I/O functions.

- **`struct VDIFile *vdiOpen(char *fn)`**

Open the file whose name is given. The filename can be either a relative or absolute path. The function returns a pointer to a dynamically created VDI file structure (see above), or a null pointer if there was an error. The function should load the header and the translation map, set the cursor to 0 and set the file descriptor to whatever was returned from the `open ()` system call.

If you're using a class, then this can return a boolean to indicate success or failure of the open.

- **void vdiClose(struct VDIFile *f)**

Close the file whose pointer is given. Deallocate any dynamically created memory regions.

- **ssize_t vdiRead(struct VDIFile *f, void *buf, size_t count)**

Reads the given number of bytes from the given VDI file's disk space, placing the bytes in the given buffer. The location of the first byte read is given by the cursor, which is relative to the start of the VDI file's data space. Advance the cursor by the number of bytes read.

- **ssize_t vdiWrite(struct VDIFile *f, void *buf, size_t count)**

Writes the given number of bytes to the given file, starting at the cursor. Bytes are written sequentially and the cursor is advanced to the end of the written block. Bytes to be written are located in the given buffer.

- **off_t vdiSeek(VDIFile *f, off_t offset, int anchor)**

Move the cursor of the given file to the given location, based on the offset and anchor values. If the resulting location is negative or larger than the disk size, do not change the value of the cursor.

If you are using a class, then the `VDIFile *` parameter is omitted.

The functions present the raw disk image as if it were the only data in the VDI file; all aspects of the VDI structure are to be hidden from the user. So, when the user wants to read, they only read from the disk space, never from the raw VDI file. Writing is similar. The cursor can only specify locations within the disk image.

If your project handles dynamic VDI files, then disk image should be presented as if all pages were laid out in the proper order, regardless of where they are in the VDI file. Missing pages must also be handled as if they were there.

You should also write a function that takes a pointer to a `VDIFile` as a parameter and displays its header fields in an easy-to-read manner. See example 1 below for a sample; your exact format may vary.

Suggestions

- `vdiSeek()` should only set the cursor in the `VDIFile` structure; it should *not* call `lseek()`.
- Reading and writing should be done one page at a time. While there are bytes left to be read or written, do the following tasks:
 1. Determine where within the current page reading or writing should begin.
 2. Determine how many bytes should be read or written in the current page.
 3. Determine the physical location of the page. This may involve page translation and/or page allocation. Questions: What if the page is not allocated? What if the page is marked as "all zeroes"?
 4. Use `lseek()` to go to the proper location within the physical page.
 5. Use `read()` or `write()` to read or write *only* the bytes within the current page.
 6. Advance the cursor by the number of bytes read or written; subtract the number of bytes read or written from the number of bytes remaining.

- If you are planning to write into the filesystem, consider using the `mmap ()` and `munmap ()` functions to load the VDI file header and translation map. These act like a write-through cache; they read the areas from the file into memory set up by the OS, and writing to them automatically writes back to the file.

Example 1

This is the output from my **testVDI()** function on the good-fixed-1k.vdi file.

Dump of VDI header:

```

Image name: [<<< Oracle VM VirtualBox Disk Image >>>
            ]
Signature: 0xbeda107f
Version: 1.01
Header size: 0x000000190 400
Image type: 0x000000002
Flags: 0x000000000
Virtual CHS: 0-0-0
Sector size: 0x000000200 512
Logical CHS: 260-16-63
Sector size: 0x000000200 512
Map offset: 0x00100000 1048576
Frame offset: 0x00200000 2097152
Frame size: 0x00100000 1048576
Extra frame size: 0x000000000 0
Total frames: 0x000000080 128
Frames allocated: 0x000000080 128
Disk size: 0x00000000008000000 134217728
UUID: 6fd7a9b3-5226-8740-aea7-506076e113b0
Last snap UUID: 6221a1ea-f266-e946-aeeb-4959370a749b
Link UUID: 00000000-0000-0000-0000-000000000000
Parent UUID: 00000000-0000-0000-0000-000000000000
Image comment:
Offset: 0x54

```

```

00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f  0...4...8...c...
+-----+-----+
00|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|00|
10|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|10|
20|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|20|
30|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|30|
40|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|40|
50|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|50|
60|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|60|
70|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|70|
80|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|80|
90|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|90|
a0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|a0|
b0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|b0|
c0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|c0|
d0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|d0|
e0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|e0|
f0|00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00|f0|
+-----+-----+

```

Unused: 0

Header in binary:

Offset: 0x0

	00	01	02	03	04	05	06	07	08	09	0a	0b	0c	0d	0e	0f	0...4...8...c...
00	3c	3c	3c	20	4f	72	61	63	6c	65	20	56	4d	20	56	69	<<< Oracle VM VirtualBox Disk Image >>>
10	72	74	75	61	6c	42	6f	78	20	44	69	73	6b	20	49	6d	
20	61	67	65	20	3e	3e	3e	0a	00	00	00	00	00	00	00	00	
30	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
40	7f	10	da	be	01	00	01	00	90	01	00	00	02	00	00	00	
50	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
60	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
70	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
80	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
90	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
a0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
b0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
c0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
d0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
e0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
f0	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	

Offset: 0x100

	00	01	02	03	04	05	06	07	08	09	0a	0b	0c	0d	0e	0f	0...4...8...c...
00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	o R& @
10	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
20	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
30	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
40	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	
50	00	00	00	00	00	00	10	00	00	00	20	00	00	00	00	00	
60	00	00	00	00	00	00	00	00	00	02	00	00	00	00	00	00	
70	00	00	00	08	00	00	00	00	00	10	00	00	00	00	00	00	
80	80	00	00	00	80	00	00	00	6f	d7	a9	b3	52	26	87	40	
90																	
a0																	
b0																	
c0																	
d0																	
e0																	
f0																	

Partition table from Master Boot Record:

Offset: 0x100

00 01 02 03 04 05 06 07 08 09 0a 0b 0c 0d 0e 0f																0...4...8...c...										
+-----+-----+																+-----+										
00																	00									
10																	10									
20																	20									
30																	30									
40																	40									
50																	50									
60																	60									
70																	70									
80																	80									
90																	90									
a0																	a0									
b0																00 04		b0								
c0		01	04	83	01	82	02	00	08	00	00	00	00	f8	03	00	00	00	00		c0					
d0		00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00		d0					
e0		00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00		e0					
f0		00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00		f0					
+-----+-----+																+-----+										

Other Examples

The output for the other five sample VDI files is available in the repository in the file `step1.log`.